

A stylized, light brown illustration of a plant with several leaves and a cluster of small, round fruits or berries, positioned on the left side of the slide.

DATABASE MANAGEMENT SYSTEM

Subash Manandhar

Chapter 5 : Relational Database Design

- **Assertions and Triggers**

- **Assertions:**

- Are general purpose checks that allow the enforcement of any condition over the entire database.
 - When the assertion is made, the system tests its validity and tests it again on every update that may violate the assertion.
 - This testing may introduce a significant amount of overhead, hence assertion should be used with great care.

Chapter 5 : Database Constraints and Relational Database Design

- **Assertions:**

- CREATE ASSERTION <assertion-name> CHECK <predicate>
- E.g. The department id of manager relation is always not null since each manager works at least one department
- CREATE ASSERTION nomanager CHECK
(NOT EXISTS (SELECT * FROM manager WHERE deptid = NULL))

It assures that there is no manager who is not assigned any departments at any time.

mid	mname	deptid
m1	Ram	d1
m2	Sita	d2
m3	Hari	NULL
Relation: manager		

Here, deptid of manager Hari is NULL, so assertion is violated and we cannot further modify database.

Chapter 5 : Database Constraints and Relational Database Design

- **Triggers:**

- is a stored procedure in database which automatically invokes whenever a special event in the database occurs. Like a trigger can be invoked when a row is inserted into a specified table or when certain table columns are being updated.

- Syntax:

CREATE TRIGGER [trigger-name] => *creates a trigger with trigger name*

[before | after] => *specifies when trigger will be executed*

[insert | update | delete] => *specifies DML operation*

ON table-name => *specifies name of table associated with trigger*

[FOR EACH ROW] => *specifies row level trigger*

[trigger-body] => *provides the operation to be performed as trigger is fired*

Chapter 5 : Database Constraints and Relational Database Design

- **Triggers:**

- Before triggers run the trigger action before the triggering statement is run.
- After trigger run the trigger action after the triggering statement is run.
- E.g. student (id,name,marks1,marks2,marks3,total,average)

```
CREATE TRIGGER stud-marks
```

```
BEFORE INSERT
```

```
ON student
```

```
FOR EACH ROW
```

```
SET new.total = new.marks1 + new.marks2 + new.marks3 ,  
    new.average = new.total/3
```


Chapter 5 : Database Constraints and Relational Database Design

- **Triggers:**

- INSERT INTO student VALUES (1,"Ram",40,50,90,0,0)

- SELECT * FROM student



1	Ram	40	50	90	180	60
---	-----	----	----	----	-----	----

Chapter 8: Transactions and Concurrency Control

- **SQL Standard Isolation Levels:**

- SQL isolation levels play a crucial role in ensuring data consistency and integrity when multiple transactions are executed concurrently.

- **Read Phenomena:**

- The standard SQL 92 defines three read phenomena describing various issues that may happen when two transactions are executed concurrently with no transaction isolation in place.

eid	name	salary
1	Ram	10000
2	Sita	15000

Fig: employee table

Chapter 8: Transactions and Concurrency Control

- **SQL Standard Isolation Levels:**

- **Read Phenomena:**

- **Dirty Read**

- When two transactions access the same data and we allow for reading values that are not yet committed, we may get a dirty read.

Transaction1	Transaction2
	UPDATE employee SET salary=15000 WHERE eid=1
SELECT salary FROM employee WHERE eid=1	
	ROLLBACK

- *Here, Transaction 2 modifies row with id = 1, then Transaction 1 reads the row and gets value 15000, and Transaction 2 rolls things back. Effectively, Transaction 1 uses value that doesn't exist in the database.*

Chapter 8: Transactions and Concurrency Control

- **SQL Standard Isolation Levels:**

- **Read Phenomena:**

- **Non Repeatable Read**

- Non Repeatable read is a problem when a transaction reads the same thing twice and gets different results each time.

Transaction1	Transaction2
SELECT salary FROM employee WHERE eid=1	
	UPDATE employee SET salary=15000 WHERE eid=1 COMMIT
SELECT salary FROM employee WHERE eid=1	

- *Transaction 1* reads a row and gets value 10000. *Transaction 2* modifies the same row. Then *Transaction 1* reads the row again and gets a different value (15000 this time).

Chapter 8: Transactions and Concurrency Control

- **SQL Standard Isolation Levels:**

- **Read Phenomena:**

- **Phantom Read**

- Phantom read is a case when a transaction looks for rows the same way twice but gets different results.

Transaction1	Transaction2
SELECT * FROM employee WHERE salary < 20000	
	INSERT INTO employee VALUES (3,"Hari",10000) COMMIT
SELECT * FROM employee WHERE salary<20000	

- *Transaction 1* reads rows and finds two of them matching the conditions. *Transaction 2* adds another row that matches the conditions used by the *Transaction 1*. When the *Transaction 1* reads again, it gets a different set of rows. We would expect to get the same rows for both *SELECT* statements of *Transaction 1*.

Chapter 8: Transactions and Concurrency Control

- **SQL Standard Isolation Levels:**

- **ISOLATION LEVELS:**

- There are 4 standard levels: *READ UNCOMMITTED*, *READ COMMITTED*, *REPEATABLE READ*, and *SERIALIZABLE*.
- ***READ UNCOMMITTED*** allows a transaction to read data that is not yet committed to the database. This allows for highest performance but it also leads to most undesired read phenomena.
- ***READ COMMITTED*** allows a transaction to read only data that is committed. This avoids issue of reading data that “later disappears” but doesn’t protect from other read phenomena.
- ***REPEATABLE READ*** level tries to avoid issue of reading data twice and getting different results.
- Finally, ***SERIALIZABLE*** tries to avoid all read phenomena.

Chapter 8: Transactions and Concurrency Control

- **SQL Standard Isolation Levels:**
 - **ISOLATION LEVELS:**

Level \ Phenomena	Dirty Read	Repeatable Read	Phantom Read
READ UNCOMMITTED	YES	YES	YES
READ COMMITTED	NO	YES	YES
REPEATABLE READ	NO	NO	YES
SERIALIZABLE	NO	NO	NO

Chapter 8 : Transactions and Concurrency Control

- **Graph based Protocol:**
 - These protocols offer an alternative to traditional locking mechanisms, often providing better performance and scalability, especially in high-contention environments.
- **Advantages of Graph-Based Protocols**
 - **Improved Concurrency:** By allowing more flexibility in the order of transaction execution, graph-based protocols can often achieve higher concurrency levels than traditional locking protocols.
 - **Reduced Lock Contention:** Since transactions do not necessarily need to acquire locks on all data items they access, lock contention can be reduced.
 - **Scalability:** Graph-based protocols can scale well to large numbers of concurrent transactions.


Chapter 8 : Transactions and Concurrency Control

- **How Graph-Based Protocols Work**

- **Transaction Submission:** When a transaction is submitted, it is assigned a timestamp.
- **Conflict Detection:** As transactions execute, conflicts are detected and recorded in the transaction graph.
- **Cycle Detection:** The graph is periodically checked for cycles. If a cycle is found, it indicates a deadlock, and one or more transactions must be aborted to break the cycle.
- **Serializability Enforcement:** The protocol ensures that the final schedule is equivalent to a serial schedule by enforcing a partial order on the transactions based on the timestamp and dependency information in the graph.

Graph Based Protocol

- Graph-based protocols are an alternative to two-phase locking
- Impose a partial ordering \rightarrow on the set $\mathbf{D} = \{d_1, d_2, \dots, d_h\}$ of all data items.
 - If $d_i \rightarrow d_j$ then any transaction accessing both d_i and d_j must access d_i before accessing d_j .
 - Implies that the set \mathbf{D} may now be viewed as a directed acyclic graph (DAG), called a *database graph*.
- The *tree-protocol* is a simple kind of graph protocol.
 - Only exclusive locks are allowed.
 - The first lock by T_i may be on any data item if there is no lock on the data item.
 - Subsequently, a data Q can be locked by T_i only if the parent of Q is currently locked by T_i .
 - Data items may be unlocked at any time.
 - in tree locking protocol a transaction may have to lock data items that it does not access.
 - increased locking overhead and additional waiting time
 - potential decrease in concurrency
 - it is deadlock free so no rollback
 - unlocking may occur earlier



	11	12	13
1	Lock-X(A)		
2	Lock-X(B)		
3		Lock-X(D)	
4		Lock-X(H)	
5		UnLock-X(D)	
6	Lock-X(E)		
7	Lock-X(D)		
8	UnLock-X(B)		
9	UnLock-X(E)		
10			Lock-X(B)
11			Lock-X(E)
12		UnLock-X(H)	
13	Lock-X(B)		
14	Lock-X(G)		
15	UnLock-X(D)		
16			UnLock-X(E)
17			UnLock-X(B)
18	UnLock-X(G)		

	I1	I2	I3
1	Lock-X(A)		
2	Lock-X(B)		
3		Lock-X(D)	
4		Lock-X(H)	
5		UnLock-X(D)	
6	Lock-X(E)		
7	Lock-X(D)		
8	UnLock-X(B)		
9	UnLock-X(E)		
10			Lock-X(B)
11			Lock-X(E)
12		UnLock-X(H)	
13	Lock-X(G)		
14	UnLock-X(D)		
15			UnLock-X(E)
16			UnLock-X(B)
17	UnLock-X(G)		