

A stylized, light brown illustration of a plant with several leaves and a cluster of small, round fruits or berries, positioned on the left side of the slide.

# DATABASE MANAGEMENT SYSTEM

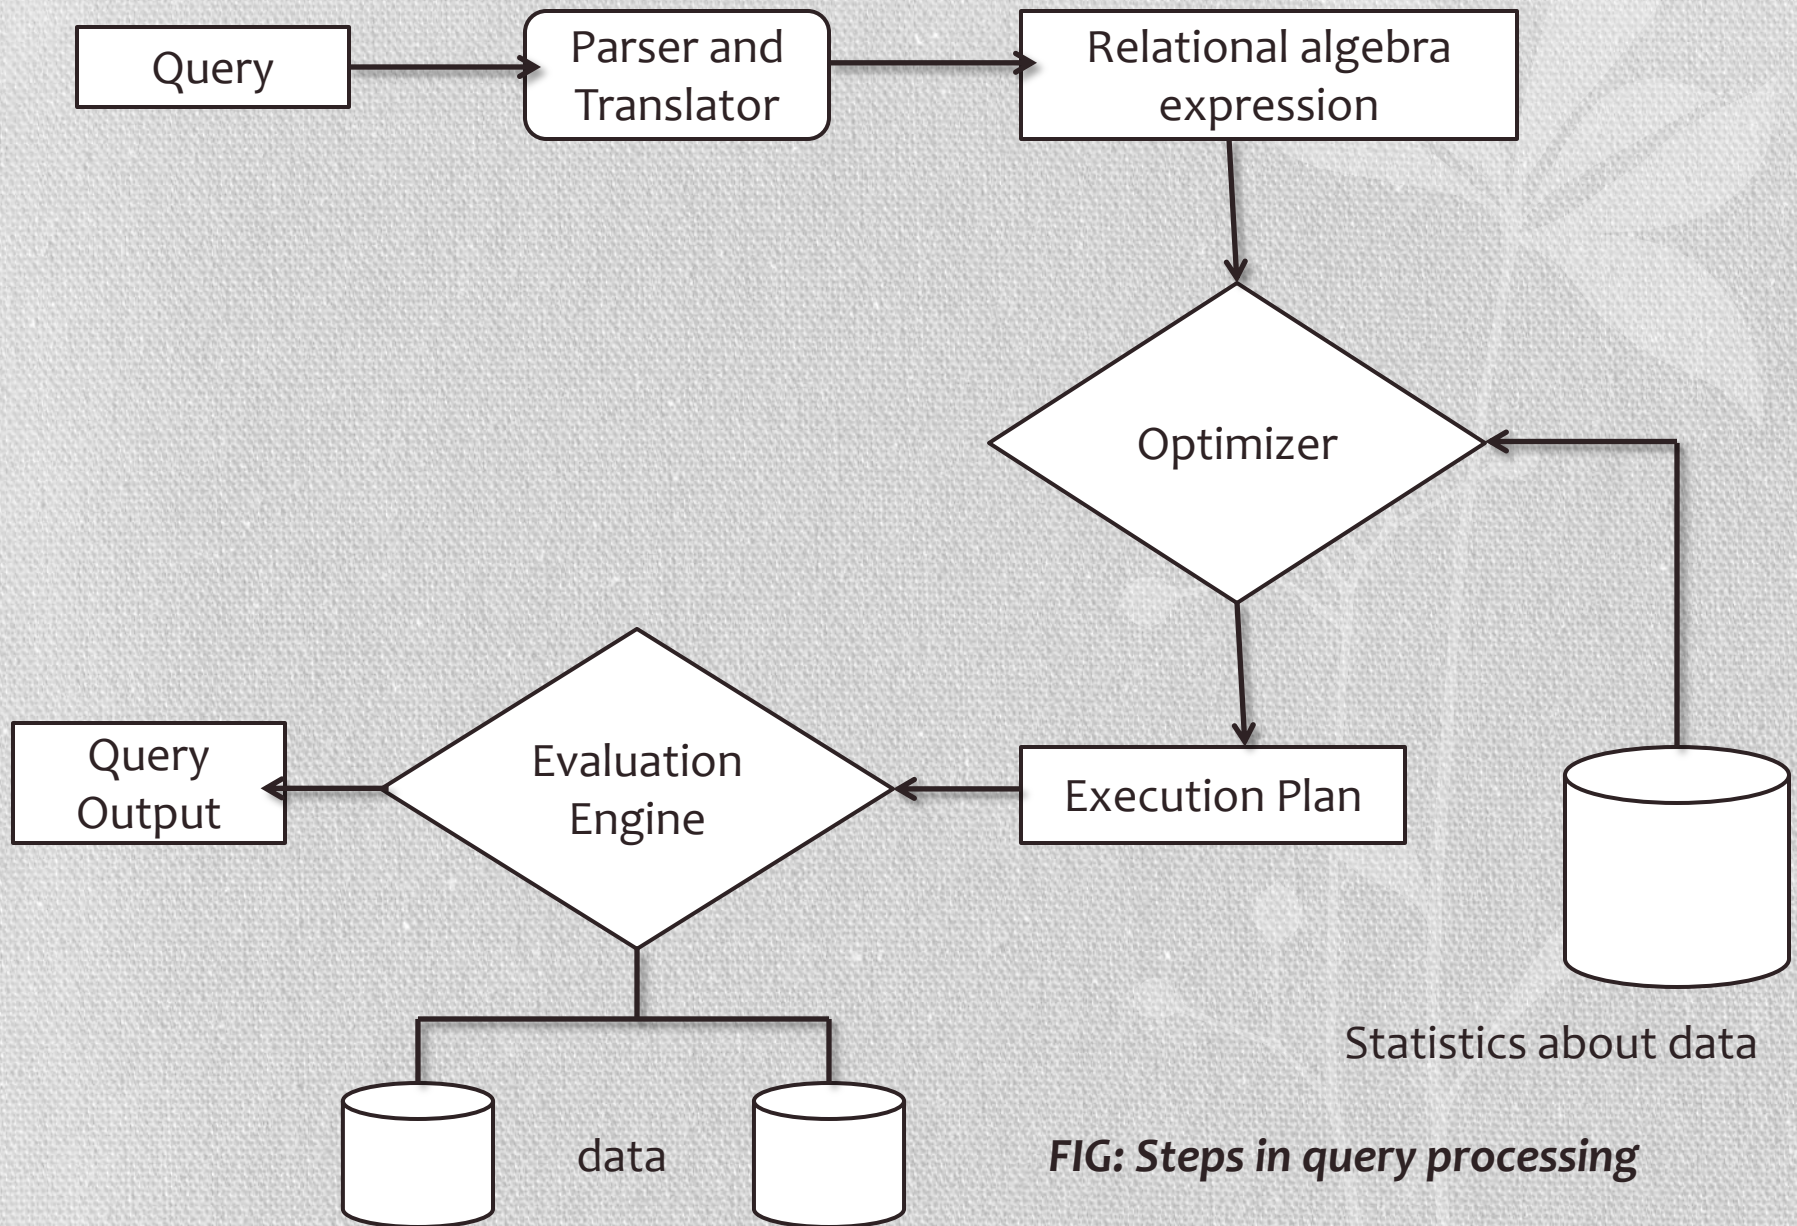
Subash Manandhar



# Chapter 6 : Query Processing and Optimization

- Query processing is transforming a query written in high level language typically in SQL into correct and efficient execution strategy expressed in a low level language and to execute the strategy to retrieve required data.
- Basic steps in query processing:
  - Parsing and translation
  - Optimization
  - Evaluation





**FIG: Steps in query processing**



# Chapter 6 : Query Processing and Optimization

- **Parsing and Translating:**

- Translate the query into internal form. This is then translated into relational algebra.
- Parser checks syntax, verifies relation.
- The relational algebra representation of a query specifies only partially how to evaluate a query.
- A relational algebra expression may have many equivalent expressions.

- Consider a query:

SELECT salary FROM employee WHERE salary > 20000

Now, equivalent relational algebra expressions are

1.  $\sigma_{\text{salary} > 20000}(\pi_{\text{salary}}(\text{employee}))$
2.  $\pi_{\text{salary}}(\sigma_{\text{salary} > 20000}(\text{employee}))$



# Chapter 6 : Query Processing and Optimization

- **Evaluation:**

- The query execution engine takes a query evaluation plan, executes that plan and returns the answers to query.

- **Query Optimization:**

- Amongst all equivalent evaluation plans choose the one with lowest cost.
- Cost is estimated using statistical information from database catalog.

e.g. no. of tuples in each relation, size of tuples etc.

A sequence of primitive operations that can be used to evaluate a query is called query execution plan or query evaluation plan.


Once the query plan is chosen, the query is evaluated with that plan and result of query is output.



# Chapter 6 : Query Processing and Optimization

- **Construction of parse tree:**

- A leaf node is created for each base relation in query.
- A non leaf node is created for each intermediate relation produced by relational algebra expression.
- The root node represents result of query.
- The sequence of operation directed from the leaves to the root.
- E.g.  $\pi_{\text{name, street}}(\sigma_{\text{status} = \text{'open'}}(\text{order} \bowtie \text{customer}))$





# Chapter 6 : Query Processing and Optimization

- **Query Optimization:**

- Is an important aspect of query processing.
- Is the activity of choosing an efficient execution strategy for processing a query.
- Aim is to choose one of many equivalent transformation that minimize resource usage.
- Reduce the total execution time of the query, which is the sum of execution times of all individual operations that make up query.
- There can be enormous differences in term of performance between different evaluation plans for same query.
- E.g. seconds Vs days to execute same query.



# Chapter 6 : Query Processing and Optimization

- **Cost based query Optimization:**

- Generate logically equivalent expressions by using a set of equivalence rules.
- Annotate the expressions to get alternative query evaluation plans.
- Select the cheapest plan based on estimated cost.
- Estimation of query evaluation cost based on statistical information from the catalog manager in combination with expected performance algorithms.

- **Importance of query Optimization:**

- Provides faster query processing
- Requires less cost per query
- Provides high performance of the system
- Consumes less memory



# Chapter 6 : Query Processing and Optimization

- **Measure of query cost:**

- Cost is generally measured as total elapsed time for answering query
  - Many factors contribute to time cost
    - *disk accesses, CPU, or even network communication*
- Typically disk access is the predominant cost, and is also relatively easy to estimate. Measured by taking into account
  - Number of seeks \* average-seek-cost
  - Number of blocks read \* average-block-read-cost
  - Number of blocks written \* average-block-write-cost
    - Cost to write a block is greater than cost to read a block
      - data is read back after being written to ensure that the write was successful
- For simplicity we just use the **number of block transfers** *from disk and the number of seeks* as the cost measures
  - $t_T$  – time to transfer one block
  - $t_S$  – time for one seek
  - Cost for  $b$  block transfers plus  $S$  seeks
$$b * t_T + S * t_S$$



# Chapter 6 : Query Processing and Optimization

- **Measure of query cost:**
  - We ignore CPU costs for simplicity
    - Real systems do take CPU cost into account
  - We do not include cost to writing output to disk in our cost formulae
- **Query operation:**
  - Selection Operation
    - The lowest level query processing operator for accessing data is the file scan.
    - search and retrieve records for a given selection condition.
  - Linear Search
    - Scan each file block and test all records to see whether they satisfy the selection condition.
    - Cost estimate =  $b_r$  block transfers + 1 seek
      - $b_r$  denotes number of blocks containing records from relation  $r$
- If selection is on a key attribute, can stop on finding record
  - cost =  $(b_r/2)$  block transfers + 1 seek
- Linear search can be applied regardless of
  - selection condition or
  - ordering of records in the file, or
  - availability of indices



# Chapter 6 : Query Processing and Optimization

- **Selection Using Indices:**

- **Index scan** – search algorithms that use an index
  - selection condition must be on search-key of index.
- **A2 (primary index, equality on key).** Retrieve a single record that satisfies the corresponding equality condition
  - $Cost = (h_i + 1) * (t_T + t_S)$
- **A3 (primary index, equality on nonkey)** Retrieve multiple records.
  - Records will be on consecutive blocks
    - Let  $b$  = number of blocks containing matching records
  - $Cost = h_i * (t_T + t_S) + t_S + t_T * b$
- **A4 (secondary index, equality on nonkey).**
  - Retrieve a single record if the search-key is a candidate key
    - $Cost = (h_i + 1) * (t_T + t_S)$
  - Retrieve multiple records if search-key is not a candidate key
    - each of  $n$  matching records may be on a different block
    - $Cost = (h_i + n) * (t_T + t_S)$ 
      - Can be very expensive!



# Chapter 6 : Query Processing and Optimization

- **Sorting:**

- A query may specify that the output should be sorted.
- The processing of some relational query operations can be implemented more efficiently based on sorted relations (join operation)
- For relations that fit into memory, techniques like quick sort can be used.
- For relations that do not fit into memory an external merge sort can be used.

- **External Merge Sort**

- Let  $M$  denote memory size (in pages).
    - **Create sorted runs.**
      - Let  $i$  be 0 initially.  
Repeatedly do the following till the end of the relation:
        - (a) Read  $M$  blocks of relation into memory
        - (b) Sort the in-memory blocks
        - (c) Write sorted data to run  $R_i$ ; increment  $i$ .
- Let the final value of  $i$  be  $N$




# Chapter 6 : Query Processing and Optimization

- **External Merge Sort**
  - **Merge the runs (N-way merge).** We assume (for now) that  $N < M$ .
    - Use  $N$  blocks of memory to buffer input runs, and 1 block to buffer output. Read the first block of each run into its buffer page
    - **repeat**
      - Select the first record (in sort order) among all buffer pages
      - Write the record to the output buffer. If the output buffer is full write it to disk.
      - Delete the record from its input buffer page.  
If the buffer page becomes empty **then**  
read the next block (if any) of the run into the buffer.
    - **until** all input buffer pages are empty:
  - If  $N \geq M$ , several merge passes are required.
    - In each pass, contiguous groups of  $M - 1$  runs are merged.
    - A pass reduces the number of runs by a factor of  $M - 1$ , and creates runs longer by the same factor.
      - E.g. If  $M=11$ , and there are 90 runs, one pass reduces the number of runs to 9, each 10 times the size of the initial runs
    - Repeated passes are performed till all runs have been merged into one.



# Chapter 6 : Query Processing and Optimization





# Chapter 6 : Query Processing and Optimization

- Equivalence Rule:

- Conjunctive selection operations can be deconstructed into a sequence of individual selections
  - $\sigma_{\theta_1 \wedge \theta_2}(E) = \sigma_{\theta_1}(\sigma_{\theta_2}(E))$
- Selection operations are *commutative*
  - $\sigma_{\theta_1}(\sigma_{\theta_2}(E)) = \sigma_{\theta_2}(\sigma_{\theta_1}(E))$
- Cascade of projection operations (only final one)
  - $\pi_{A_1}(\pi_{A_2}(\dots(\pi_{A_n}(E))\dots)) = \pi_{A_1}(E)$
- Selections can be combined with cartesian products and theta joins
  - $\sigma_{\theta}(E_1 \times E_2) = E_1 \bowtie_{\theta} E_2$
  - $\sigma_{\theta_1}(E_1 \bowtie_{\theta_2} E_2) = E_1 \bowtie_{\theta_1 \wedge \theta_2} E_2$
- Theta join (and natural join) operations are *commutative*
  - $E_1 \bowtie_{\theta} E_2 = E_2 \bowtie_{\theta} E_1$
  - note that the order of attributes is ignored
- Natural join operations are *associative*
  - $(E_1 \bowtie E_2) \bowtie E_3 = E_1 \bowtie (E_2 \bowtie E_3)$
- Theta joins are *associative* in the following manner
  - $(E_1 \bowtie_{\theta_1} E_2) \bowtie_{\theta_2 \wedge \theta_3} E_3 = E_1 \bowtie_{\theta_1 \wedge \theta_3} (E_2 \bowtie_{\theta_2} E_3)$
  - where  $\theta_2$  contains attributes only from  $E_2$  and  $E_3$
- Union and intersection operations are *commutative*
  - $E_1 \cup E_2 = E_2 \cup E_1$
  - $E_1 \cap E_2 = E_2 \cap E_1$



# Chapter 6 : Query Processing and Optimization

- Equivalence Rule:

Union and intersection operations are *associative*

- $(E_1 \cup E_2) \cup E_3 = E_1 \cup (E_2 \cup E_3)$
- $(E_1 \cap E_2) \cap E_3 = E_1 \cap (E_2 \cap E_3)$

The *selection* operation *distributes* over *union*, *intersection* and *set difference*

- $\sigma_P(E_1 - E_2) = \sigma_P(E_1) - \sigma_P(E_2)$

The *projection distributes* over the *union* operation

- $\pi_A(E_1 \cup E_2) = (\pi_A(E_1)) \cup (\pi_A(E_2))$

Note that this is only a selection of equivalence rules



# Chapter 6 : Query Processing and Optimization

- **Heuristic Optimization:**

- Cost-based optimization is expensive, even with dynamic programming.
- Systems may use *heuristics* to reduce the number of choices that must be made in a cost-based fashion.
- Heuristic optimization transforms the query-tree by using a set of rules that typically (but not in all cases) improve execution performance:
  - Perform selection early (reduces the number of tuples)
  - Perform projection early (reduces the number of attributes)
  - Perform most restrictive selection and join operations before other similar operations.
  - Some systems use only heuristics, others combine heuristics with partial cost-based optimization.



# Chapter 6 : Query Processing and Optimization

- **Steps in Heuristic Optimization:**

- Deconstruct conjunctive selections into a sequence of single selection operations .
- Move selection operations down the query tree for the earliest possible execution .
- Execute first those selection and join operations that will produce the smallest relations .
- Replace Cartesian product operations that are followed by a selection condition by join operations .
- Deconstruct and move as far down the tree as possible lists of projection attributes, creating new projections where needed .
- Identify those sub trees whose operations can be pipelined, and execute them using pipelining.



# Chapter 6 : Query Processing and Optimization

- **Performance Tuning**

- **Database Statistics**

- Get the correct and updated statistics

- **Create Optimized Indexes**

- Have a right balance of indexes on tables

- **Specify optimizer hints in SELECT**

- specify the index name in SELECT query

- **Predetermine expected growth**

- specify an appropriate value for fill factor when creating indexes

- **Select limited data**

- Rather than filtering on the client, push as much filtering as possible on the server-end
    - Eliminate any obvious or computed columns

- **Drop indexes before loading data**

- drop the indexes on a table before loading a large batch of data. This makes the insert statement run faster.
    - Recreate the indexes after insertion.

- **Avoid foreign key constraints**

- Foreign keys constraints ensure data integrity at the cost of performance